

Using Custom Code, Models and Containers in SageMaker



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Develop custom models using code in Apache MXNet and TensorFlow

Host pre-trained model artifacts on SageMaker containers

Train custom models on your own containers and host them

Use other Amazon services such as Redshift from within SageMaker

Custom Models in Apache MXNet

Apache MXNet is an effort undergoing incubation at The Apache Software Foundation (ASF), **sponsored by the *Apache Incubator***. Incubation is required of all newly accepted projects until a further review indicates that the infrastructure, communications, and decision making process have stabilized in a manner consistent with other successful ASF projects. While incubation status is not necessarily a reflection of the completeness or stability of the code, it does indicate that the project has yet to be fully endorsed by the ASF.

TensorFlow

Open source deep
learning framework by
Google

Apache MXNet

Deep learning framework
sponsored by the Apache
Incubator

Apache Spark

Engine for big data
processing with powerful
ML libraries

Custom Code in SageMaker

TensorFlow

Open source deep learning framework by Google

Apache MXNet

Deep learning framework sponsored by the Apache Incubator

Apache Spark

Engine for big data processing with powerful ML libraries

To train and host custom code on SageMaker the code needs to follow a certain **training** and **inference interface**

`train(...)`

Apache MXNet Training Code Interface

SageMaker calls this function with information on the training environment to run training

May return the model object which is passed to the `save()` function

`save(...)`

Apache MXNet Training Code Interface

Saves a model after training

This is an optional function, called only if you have a model to save

`model_fn(...)`

Apache MXNet Inference Code Interface

Loads a model from disk

```
transform_fn(...)
```

Apache MXNet Inference Code Interface

Transforms input data into a prediction result

The Apache MXNet inference interface for Gluon models (an imperative interface) are different

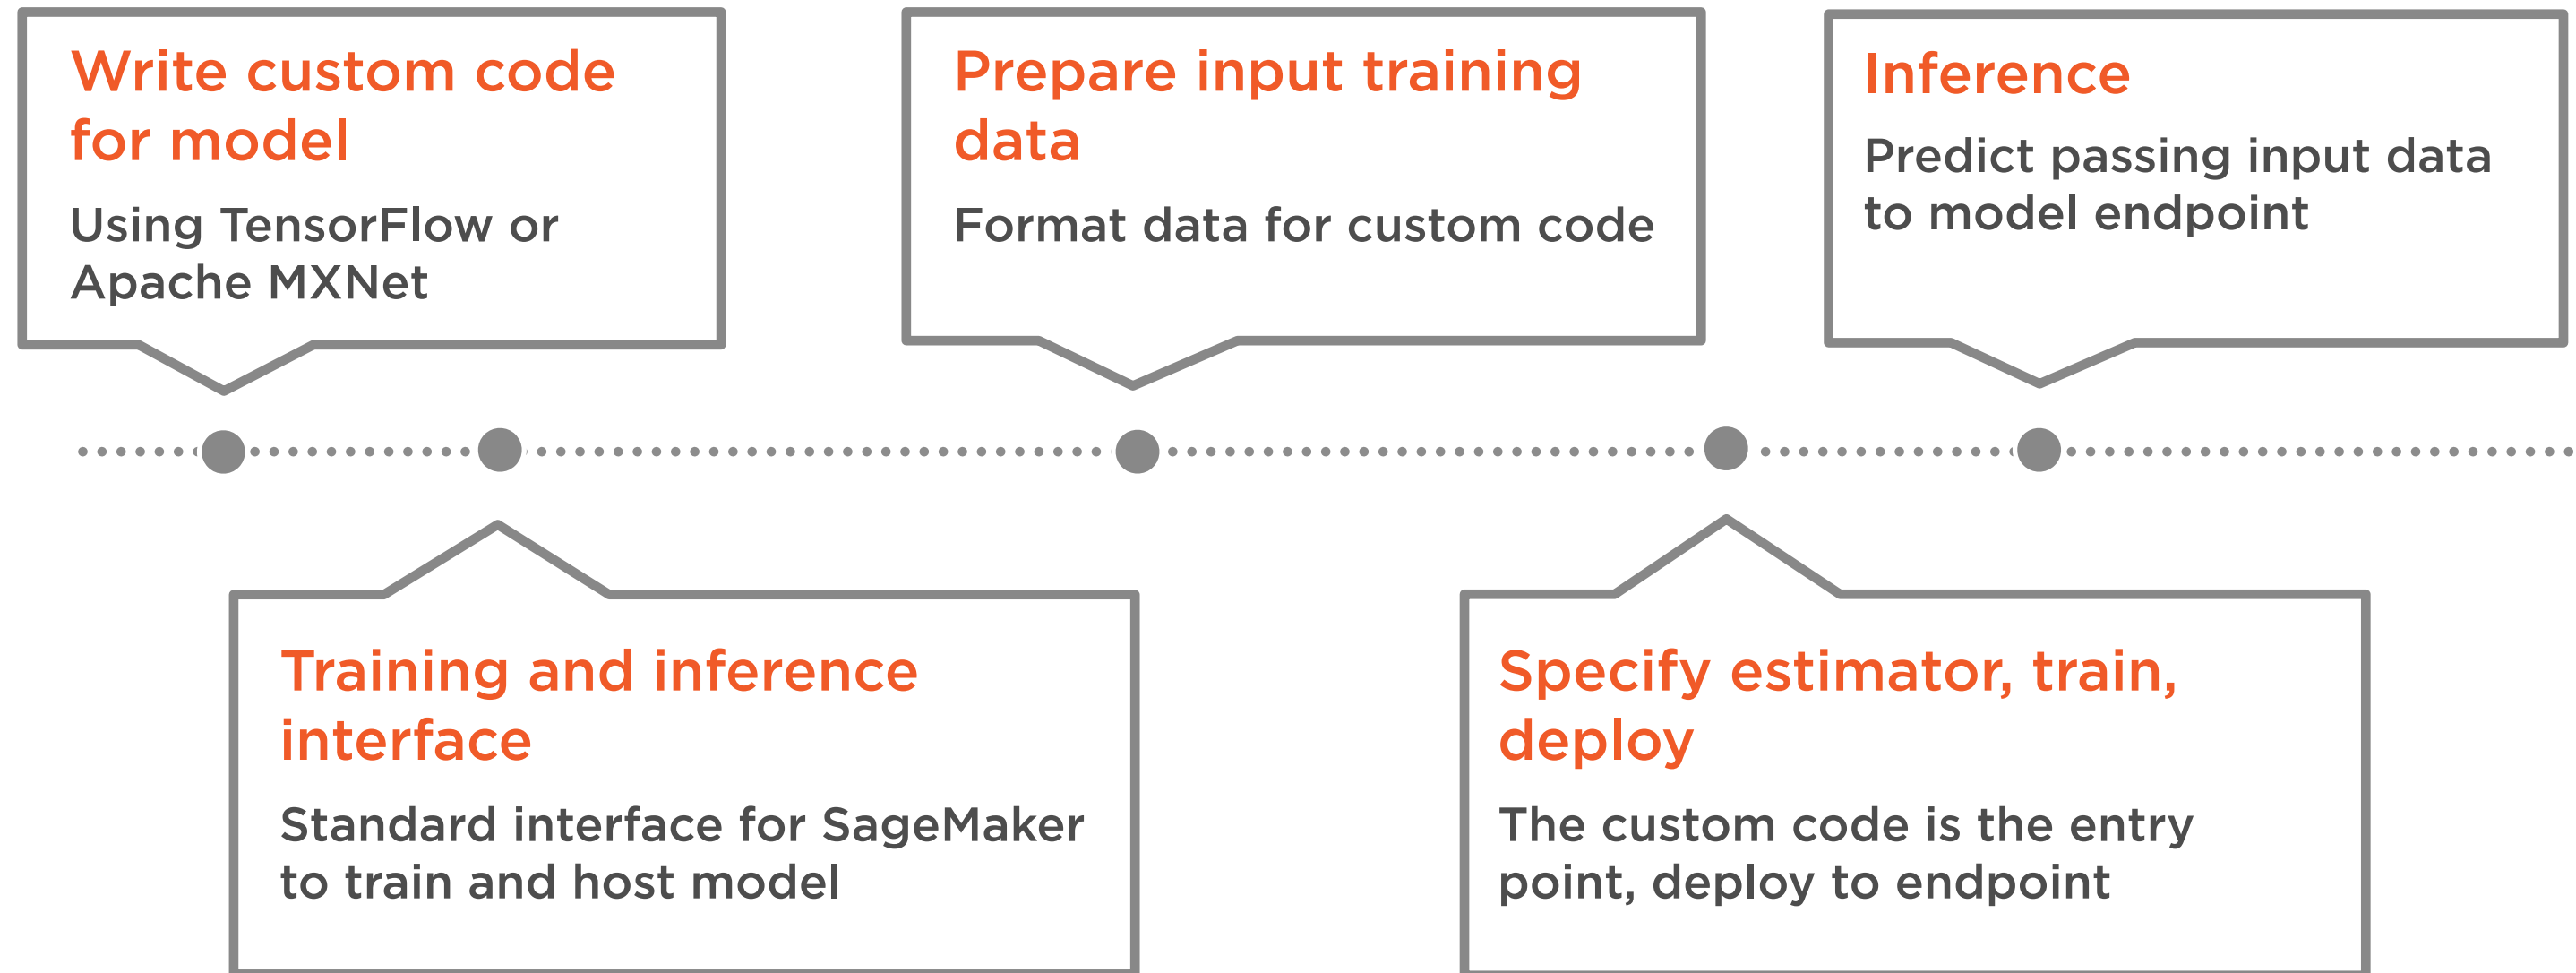
Demo

Use a neural network built in Apache MXNet for handwritten digit classification

Train on the MNIST dataset

Identify handwritten digits written on an HTML canvas

Custom Code on SageMaker



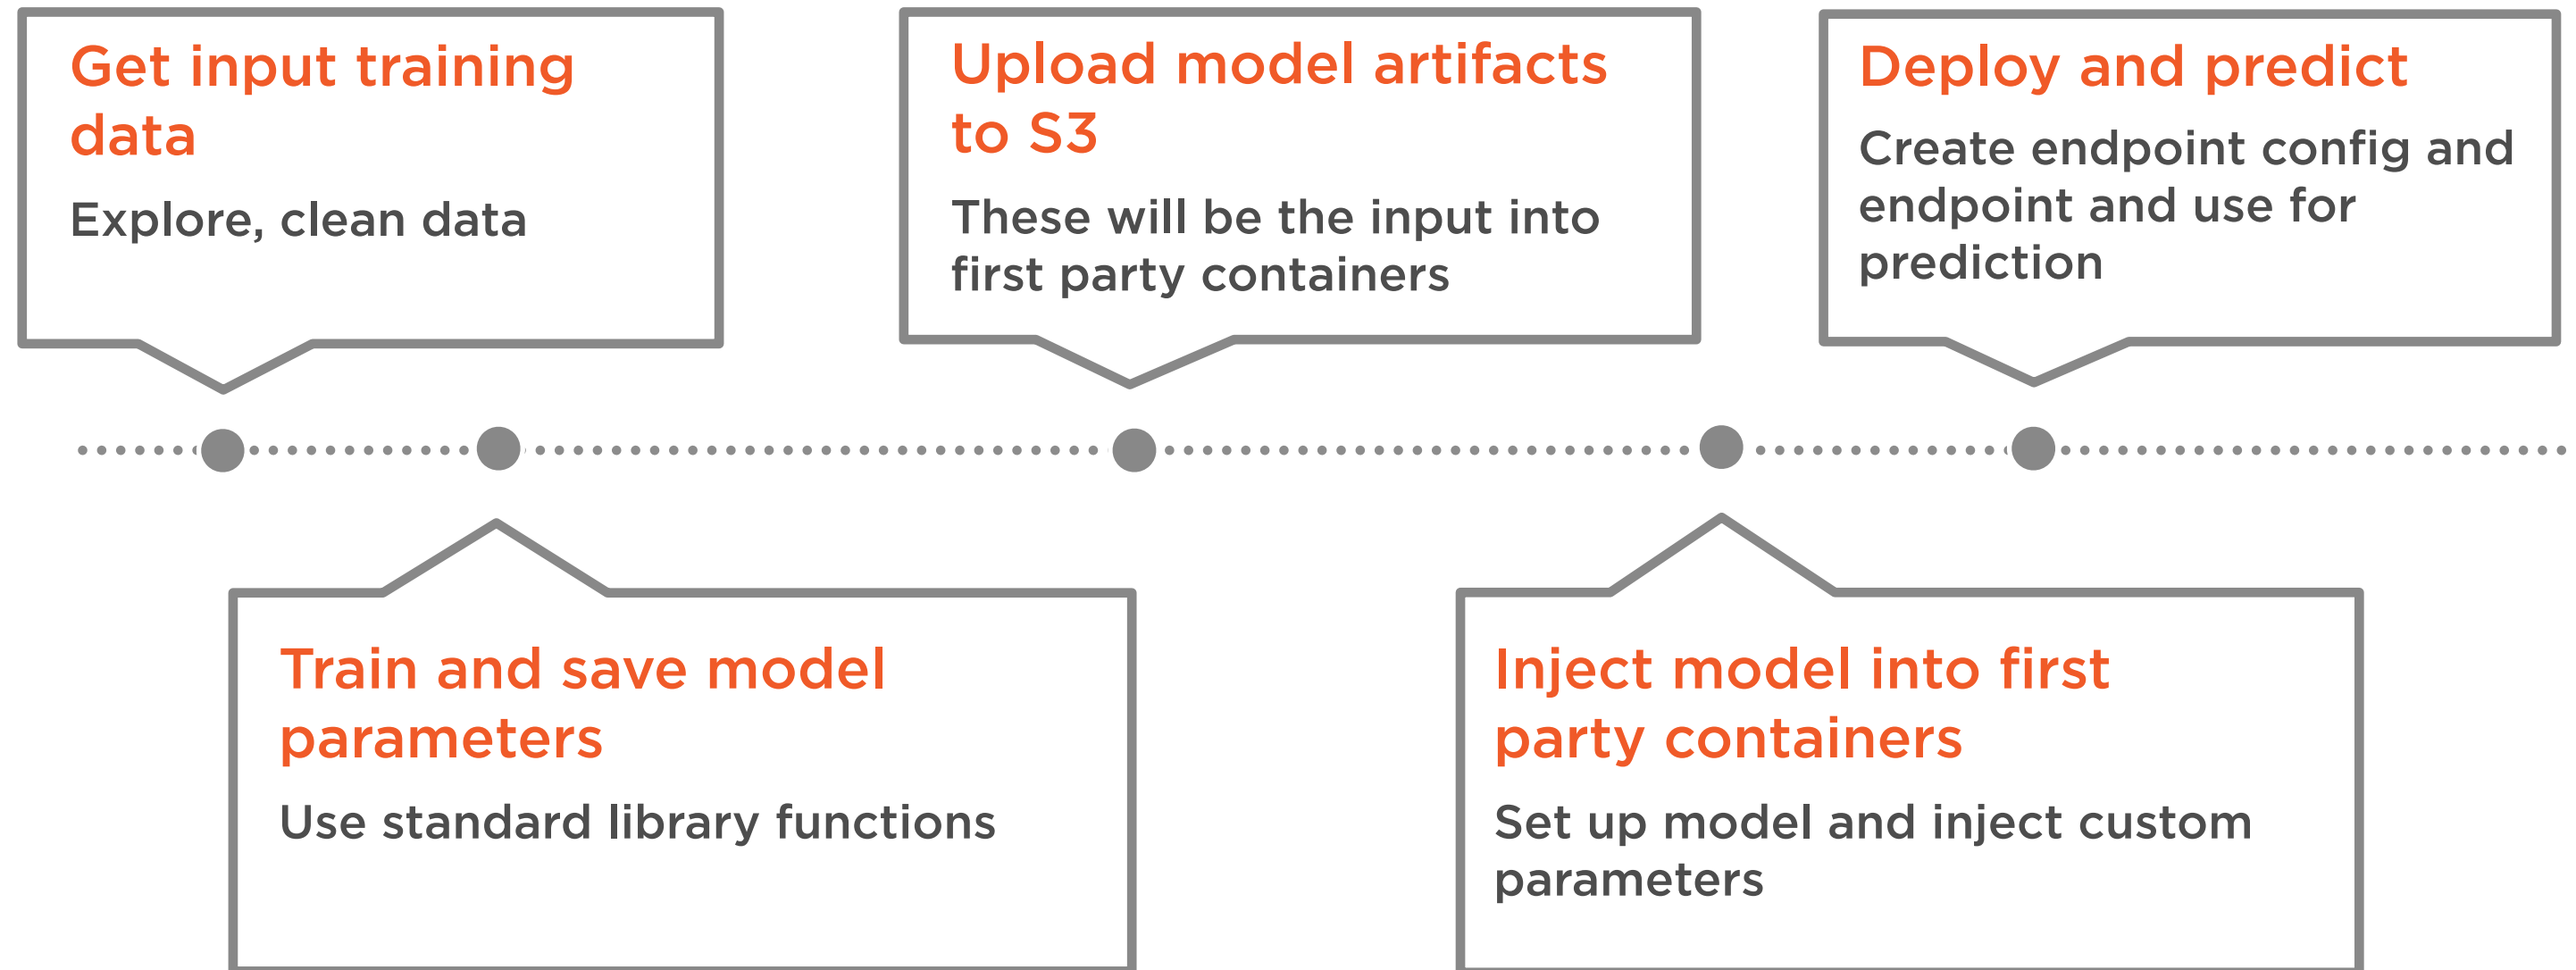
Demo

Set up the built-in algorithms to use your pre-trained model parameters

Cluster the MNIST handwritten digits using the k-means clustering algorithms in scikit-learn

Inject the cluster centers into SageMaker's first party containers

Custom Model on SageMaker



Demo

Setting up a Redshift cluster on AWS

Demo

**Connecting to and querying Redshift
from within SageMaker's Jupyter
notebook instance**

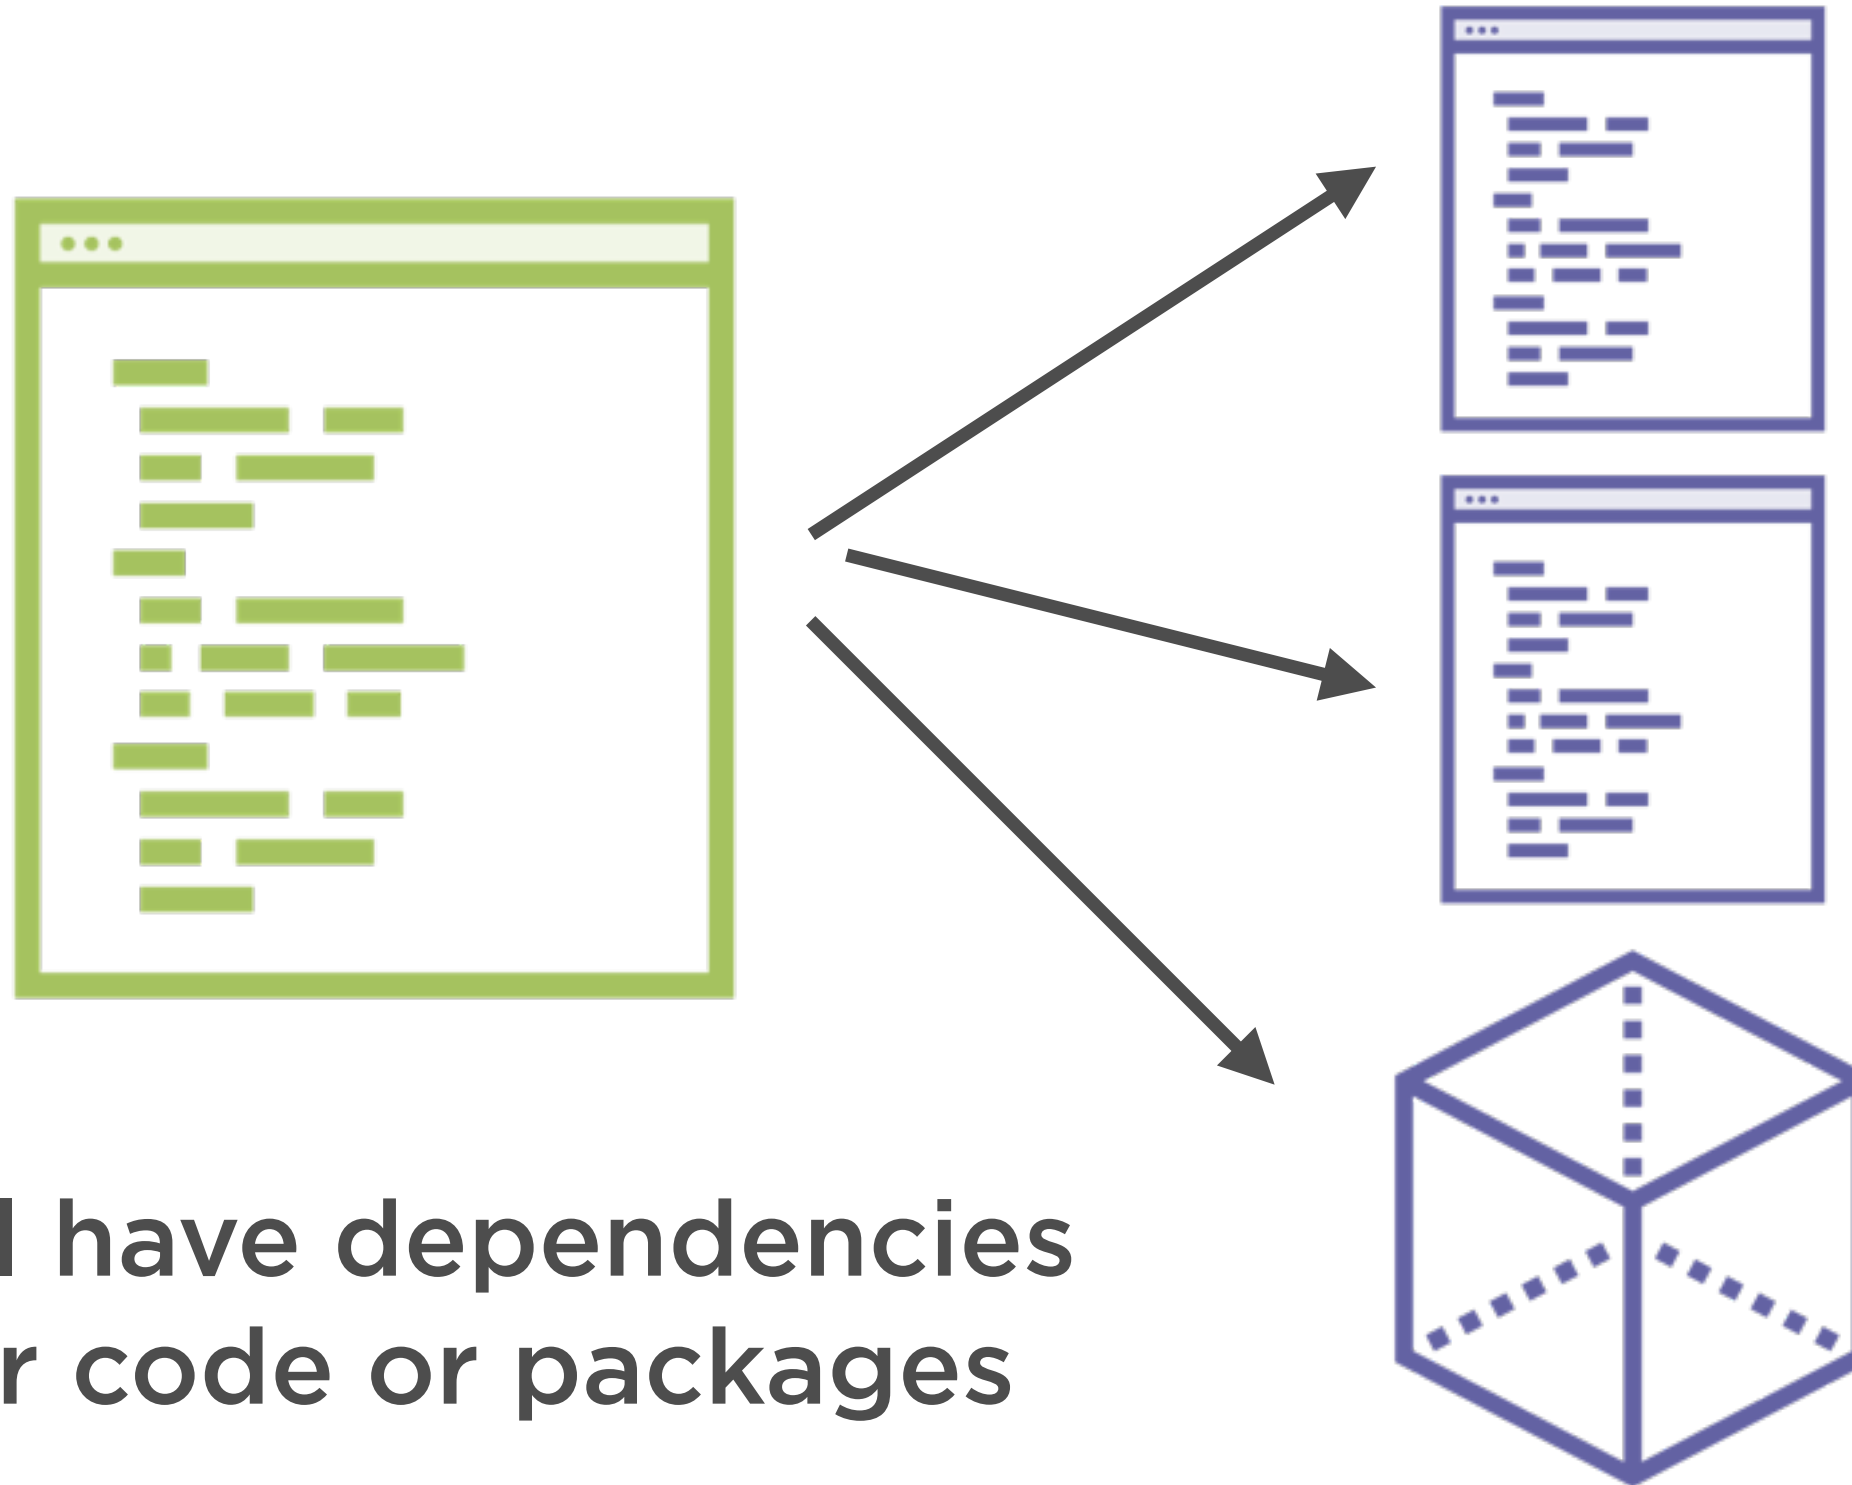
Bring Your Own Container

Custom Algorithms on SageMaker



**Custom code to train and
host our model**

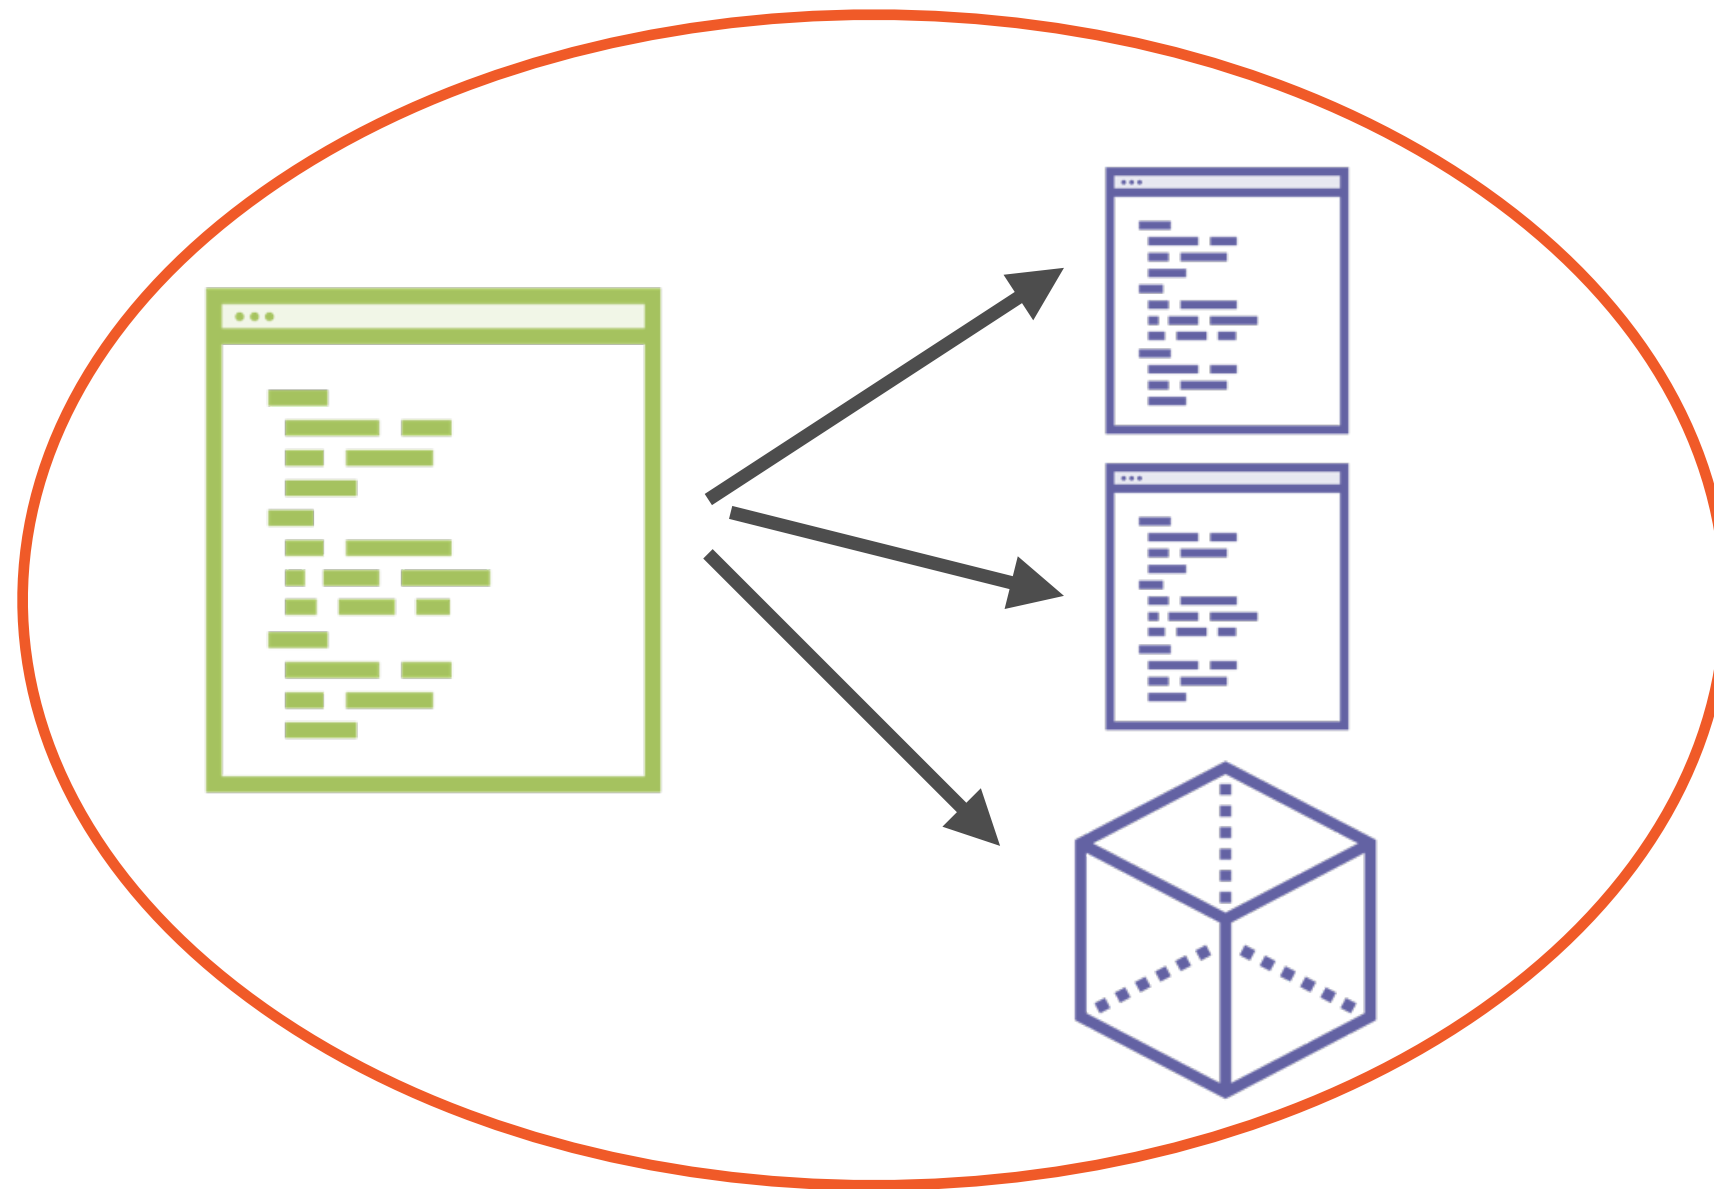
Custom Algorithms on SageMaker



**Code will have dependencies
on other code or packages**




Docker Containers



Docker packages all of this arbitrary code into an **image**




Docker Containers



**An image is completely
self-sufficient**




Docker Containers



An image can be used to run a Docker container




Docker Containers



Containers are a fully self contained environment which executes code




Docker Containers



Containers do not contain the OS and your code is abstracted away from the machine



Docker Containers



You can create as many containers as you want from the same image


[What is Docker?](#)[Product](#)[Community](#)[Support ▾](#)[Create Docker ID](#)[Sign In](#)

DOCKER PLATFORM ADDS KUBERNETES

Simplify and advance the management of Kubernetes for enterprise IT

[LEARN MORE](#)[SIGN UP FOR THE BETA](#)


SageMaker can train and host
custom ML code in Docker
containers



Getting your own container set up
on SageMaker is fairly complex

SageMaker has an end-to-
end example which you can
tweak to add in your own
custom model

Custom Algorithm



Decision tree classifier from scikit-learn to classify Iris flowers

Input features: Petal and sepal length and width

Output labels: Iris Versicolor, Iris Setosa, Iris Virginica

Running the Docker Container

Training

Run the train script and use `/opt/ml` directory to store data

Hosting

Use `nginx`, `gunicorn`, `flask` to respond to HTTP requests

Running the Docker Container

Training

Run the train script and use `/opt/ml` directory to store data

Hosting

Use `nginx`, `gunicorn`, `flask` to respond to HTTP requests

Training

`/opt/ml`

```
|— input
|   |— config
|   |   |— hyperparameters.json
|   |   |— resourceConfig.json
|   |— data
|   |   |— <channel_name>
|   |   |— <input data>
|— model
|   |— <model files>
|— output
|   |— failure
```

High level directory where
SageMaker stores information
used during the training run

Training

```
/opt/ml
├─ input
│   ├── config
│   │   └─ hyperparameters.json
│   │   └─ resourceConfig.json
│   └─ data
│       ├── <channel_name>
│       └─ <input data>
├─ model
│   └─ <model files>
└─ output
    └─ failure
```

Hyperparameters control how
your program runs

Training

```
/opt/ml
├─ input
│   ├─ config
│   │   └─ hyperparameters.json
│   │   └─ resourceConfig.json
│   └─ data
│       └─ <channel_name>
│           └─ <input data>
├─ model
│   └─ <model files>
└─ output
    └─ failure
```

Specify the network layout for distributed training

Training

```
/opt/ml
├─ input
│   └─ config
│       ├── hyperparameters.json
│       └─ resourceConfig.json
│   └─ data
│       ├── <channel_name>
│       └── <input data>
├─ model
│   └─ <model files>
└─ output
    └─ failure
```

scikit-learn libraries cannot be run
in a distributed manner - we won't
be using both of these

Training

```
/opt/ml
├─ input
│   ├── config
│   │   ├── hyperparameters.json
│   │   └── resourceConfig.json
│   └─ data
│       ├── <channel_name>
│       │   └── <input data>
├─ model
│   └── <model files>
└─ output
    └── failure
```

The input data is copied from the S3 bucket to this location

Training

```
/opt/ml
├─ input
│   ├── config
│   │   ├── hyperparameters.json
│   │   └─ resourceConfig.json
│   └─ data
│       ├── <channel_name>
│       └─ <input data>
├─ model
│   └─ <model files>
└─ output
    └─ failure
```

Saved model parameters which
are then uploaded to S3

Training

```
/opt/ml
├─ input
│   ├── config
│   │   ├── hyperparameters.json
│   │   └─ resourceConfig.json
│   └─ data
│       └─ <channel_name>
│           └─ <input data>
├─ model
│   └─ <model files>
└─ output
    └─ failure
```

Written out only in the case of
failed jobs

Running the Docker Container

Training

Run the train script and use /opt/ml directory to store data

Hosting

Use nginx, gunicorn, flask to respond to HTTP requests

Hosting

HTTP Request





nginx

nginx

Open source software for web serving, reverse proxying, caching, load balancing

Reverse proxy:

- Sits behind a firewall and directs requests to the appropriate backend
- Additional level of abstraction between client and server

Hosting

HTTP Request



gunicorn



gunicorn

Web server for Unix

WSGI HTTP Server:

- WSGI (Web Server Gateway Interface) is a Python standard which determines how a web server communicates with applications
- Simple, lightweight, fast and works with many web frameworks

Hosting

HTTP Request



nginx



gunicorn



flask



flask

flask

Microframework for Python web app development

Worker:

- The actual instance of the application which hosts the inference code
- Loads the trained model and returns prediction results

Hosting

HTTP Request



nginx



gunicorn



flask

Custom Container Components

Dockerfile

Docker command to build and run a custom container

nginx.conf

Configuration file for the reverse proxy

train and serve

Entry point for training and serving the model

build_and_push.sh

Script to create a Docker image and push to the Amazon ECR

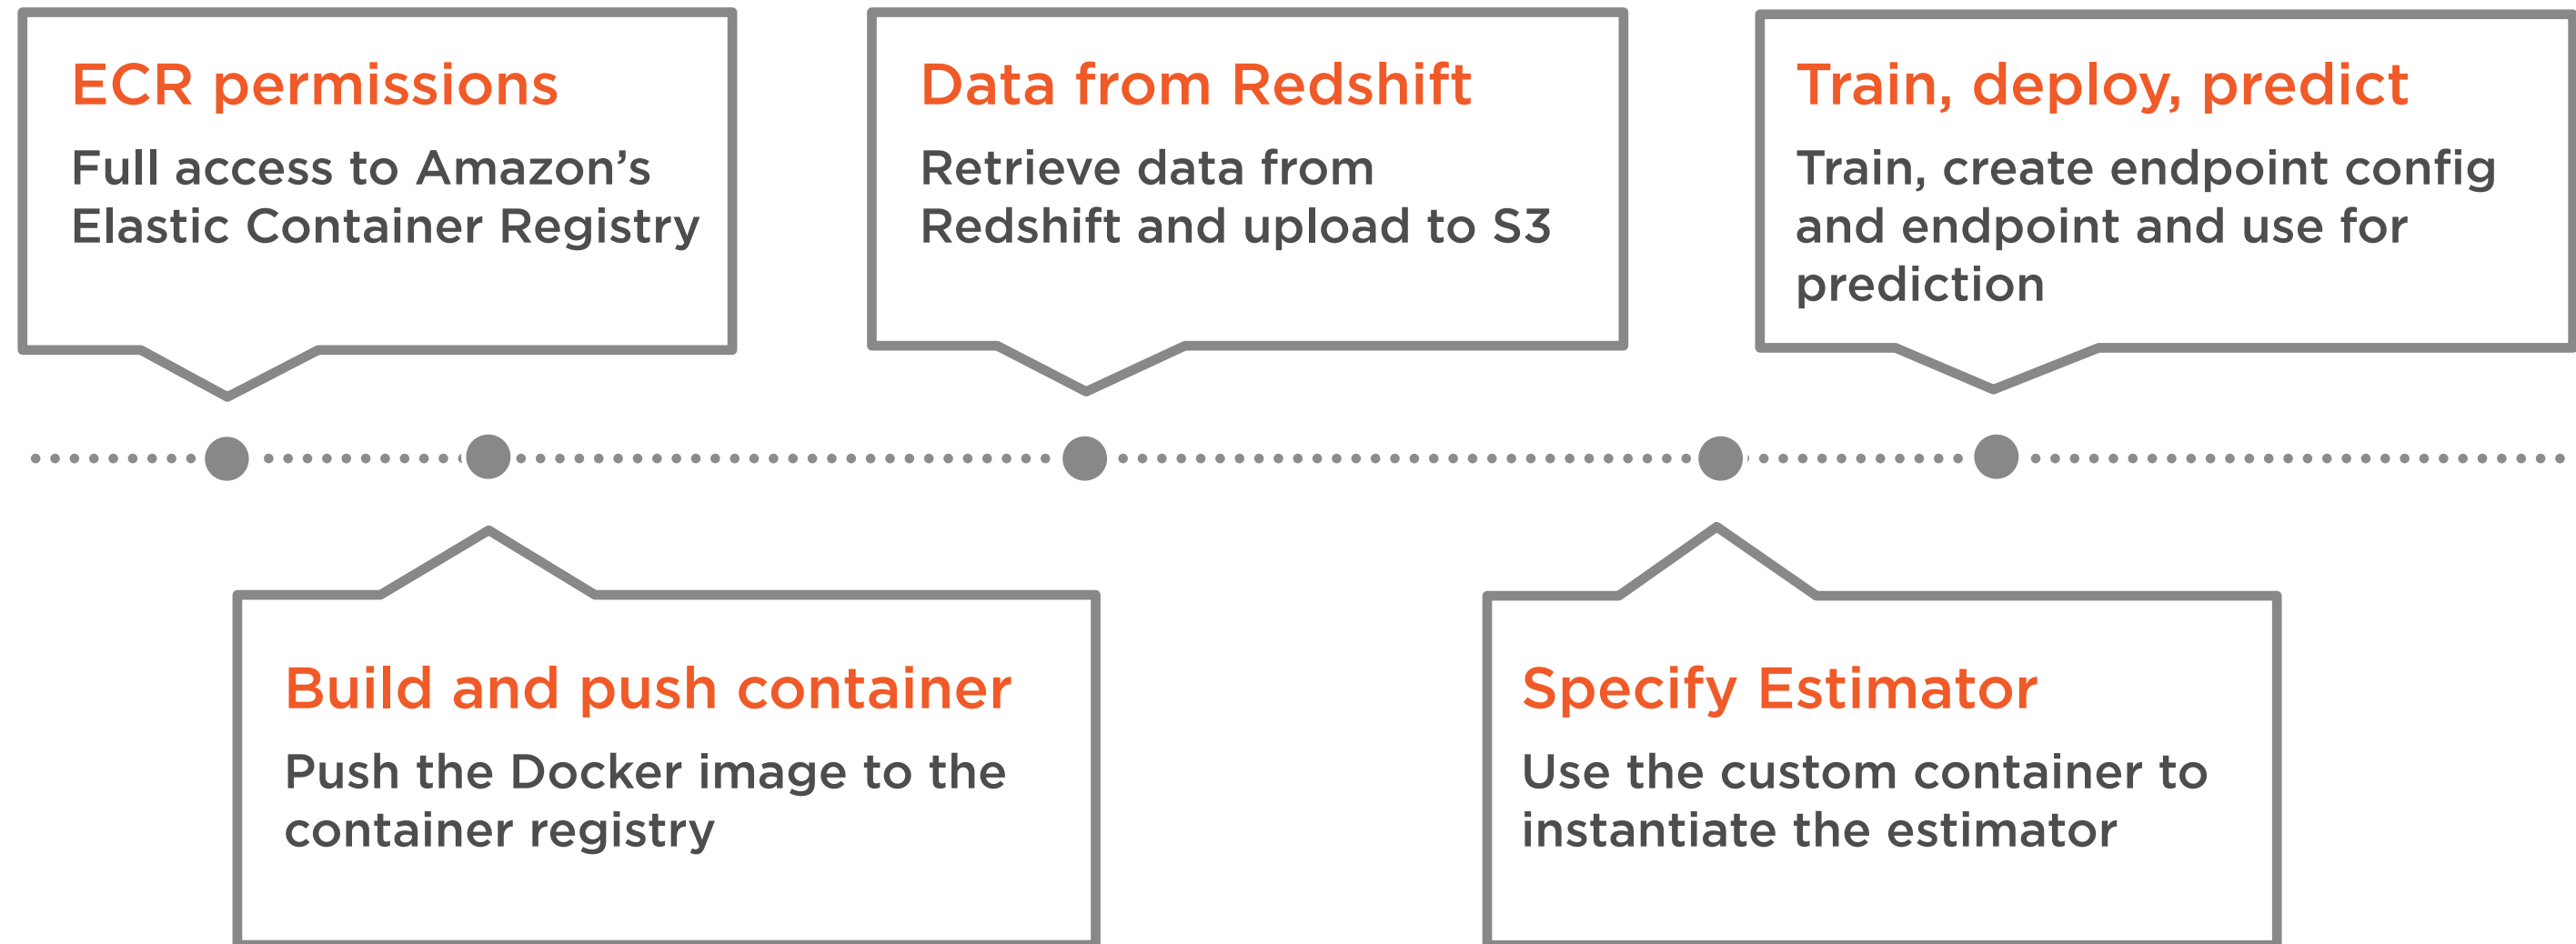
predictor.py

Flask web app for model inference

wsgi.py

Wrapper used to invoke the Flask app

Custom Container on SageMaker



Demo

Use a Docker container with custom code for training and hosting

Access data from Redshift and store to S3 buckets

Use decision trees in scikit-learn for Iris dataset classification

Summary

Develop custom models using code in Apache MXNet and TensorFlow

Host pre-trained model artifacts on SageMaker containers

Train custom models on your own containers and host them

Use other Amazon services such as Redshift from within SageMaker